

LLM API

Ignacio Gómez
2025

API Keys

K Kerox Technology / Iteldoc

Playground Dashboard Docs API reference I

DASHBOARD

Logs

Traces

Assistants

Batches

Evaluations

Fine-tuning

Storage

Usage

API keys

API keys

+ Create new secret key

As an owner of this project, you can view and manage all API keys in this project.

Do not share your API key with others or expose it in the browser or other client-side code. To protect your account's security, OpenAI may automatically disable any API key that has leaked publicly.

View usage per API key on the [Usage page](#).

NAME	SECRET KEY	CREATED	LAST USED ⓘ	CREATED BY	PERMISSIONS	
Desarrollo	sk-...lgQS	7 may 2024	25 mar 2025	Kerox Tech	All	 
Iteldoc PRO	sk-...6X1d	7 may 2024	24 mar 2025	Kerox Tech	All	 

Groq

OpenAI

Text generation and promp

Que es un LLM API | Qwen

Products - Groq is Fast AI

GroqCloud

Playground - GroqCloud

console.groq.com/docs/rate-limits

NoticiasUPMMatematicasControlRecordatorioDeep LearningPopularesSaaSGoogle

Todos los marcadores

groqcloud

PlaygroundAPI KeysDashboardDocumentationUpgradePersonal

DocsAPI Reference

GET STARTED

Overview

Quickstart

OpenAI Compatibility

Models

Rate Limits

FEATURES

Text

Speech to Text

Reasoning

Vision

ADVANCED FEATURES

Batch Processing

Flex Processing

Content Moderation

Prefilling

Tool Use

DEVELOPER RESOURCES

Rate Limits

The following is a high level summary and there may be exceptions to these limits. You can view the current, exact rate limits for your organization on the [limits page](#) in your account settings.

Free Tier

Developer Tier

MODEL ID	RPM	RPD	TPM	TPD	ASH	ASD
deepseek-r1-distill-llama-70b	30	1,000	6,000	-	-	-
deepseek-r1-distill-qwen-32b	30	1,000	6,000	-	-	-
distil-whisper-large-v3-en	20	2,000	-	-	7,200	28,800
gemma2-9b-it	30	14,400	15,000	500,000	-	-
llama-3.1-8b-instant	30	14,400	6,000	500,000	-	-
llama-3.2-1b-preview	30	7,000	7,000	500,000	-	-
llama-3.2-3b-preview	30	7,000	7,000	500,000	-	-
llama-3.2-11b-vision-preview	30	7,000	7,000	500,000	-	-
llama-3.2-90b-vision-preview	15	3,500	7,000	250,000	-	-
llama-3.3-70b-specdec	30	1,000	6,000	100,000	-	-
llama-3.3-70b-versatile	30	1,000	6,000	100,000	-	-
llama-guard-3-8b	30	14,400	15,000	500,000	-	-
llama3-8b-8192	30	14,400	6,000	500,000	-	-

Modelos de transcripción

gpt-4o-mini-transcribe

gpt-4o-transcribe

whisper

```
from openai import OpenAI

client = OpenAI()
audio_file= open("/path/to/file/audio.mp3", "rb")


transcription = client.audio.transcriptions.create(
    model="gpt-4o-transcribe",
    file=audio_file
)

print(transcription.text)
```



LangChain

LangChain es una biblioteca de código abierto diseñada para facilitar la creación de aplicaciones que utilizan modelos de lenguaje (LLMs). Permite integrar, de manera modular, diversas operaciones como la generación y manejo de prompts, la gestión de contexto y memoria, y la conexión con datos externos o herramientas adicionales.



Fine-tuning

 Copy page

Fine-tune models for better results and efficiency.

Fine-tuning lets you get more out of the models available through the API by providing:

- Higher quality results than prompting
- Ability to train on more examples than can fit in a prompt
- Token savings due to shorter prompts
- Lower latency requests

OpenAI's text generation models have been pre-trained on a vast amount of text. To use the models effectively, we include instructions and sometimes several examples in a prompt. Using demonstrations to show how to perform a task is often called "few-shot learning."

Fine-tuning improves on few-shot learning by training on many more examples than can fit in the prompt, letting you achieve better results on a wide number of tasks. **Once a model has been fine-tuned, you won't need to provide as many examples in the prompt.** This saves costs and enables lower-latency requests.

At a high level, fine-tuning involves the following steps:

- 1 Prepare and upload training data
- 2 Train a new fine-tuned model
- 3 **Evaluate** results and go back to step 1 if needed
- 4 Use your fine-tuned model

When to use fine-tuning

[Common use cases](#)

[Preparing your dataset](#)

[Create a fine-tuned model](#)

[Use a fine-tuned model](#)

[Use a checkpointed model](#)

[Analyzing your fine-tuned model](#)


[Vision fine-tuning](#)

[Preference fine-tuning](#)


[Fine-tuning examples](#)

[FAQ](#)


Retrieval Augmented Generation(RAG)



Advanced RAG with Hybrid Search



Re-Ranking



 New: Introducing the `lmstudio-python` and `lmstudio-js` SDK libraries!



Your local AI toolkit.

Download and run Llama, DeepSeek, Mistral, Phi on your computer.

 [Download LM Studio for Windows](#) 0.3.13

By using LM Studio, you agree to its [terms of use](#).

